

Maestría en Inteligencia Artificial Aplicada

Curso: Navegación autónoma

Tecnológico de Monterrey

Dr. David Antonio Torres

Actividad 4.2 - Detección de Señales de Tránsito

Nombres y matrículas de los integrantes del equipo:

- Julio Cesar Lynn Jimenez A01793660
 - Francisco Javier Parga García A01794380
 - Carlos Roberto Torres Ferguson A01215432
 - Fernando Sebastian Sanchez Cardona A01687530
-

Clasificación de señales de transito con CNN

Mediante un modelo convolucional, se creó u clasificador para identificar el tipo de señalización, utilizando los datos para entrenamiento de GTRSB

```
In [1]: import os
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Conv2D, MaxPool2D, Flatten, Dense, Dropout, BatchNormalization, ReLU
from tensorflow.keras.models import Sequential
```

```

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint
from tensorflow.keras.models import load_model
import pandas as pd
import numpy as np
import tensorflow as tf

import matplotlib.pyplot as plt
from matplotlib import rcParams
rcParams['figure.figsize'] = (5,3)

```

Datos para entrenamiento y validación

```

In [2]: train_directory = 'data/Train'
        test_directory = 'data/Test'

```

```

In [3]: # Crea un generador con aumento de datos
        train_datagen = ImageDataGenerator(rescale=1./255,
                                           horizontal_flip=True,
                                           vertical_flip=False,
                                           fill_mode='nearest',
                                           validation_split=0.25)

```

```

In [4]: # Definicion de variables para batch y tamaño de imagen a usar
        batch_size = 128
        image_size = (28,28)

        # Keras ImageDataGenerator proporciona un método flow_from_directory para cargar imágenes en lotes
        # Esto es útil para conjuntos de datos grandes que no caben en la memoria
        train_generator = train_datagen.flow_from_directory(train_directory,
                                                            target_size=image_size,
                                                            batch_size=batch_size,
                                                            class_mode='categorical',
                                                            classes= list(map(str,range(0,43))),
                                                            subset='training') # establece como datos de entrenamiento

        validation_generator = train_datagen.flow_from_directory(train_directory, # mismo directorio que los datos de entrenamiento
                                                                target_size=image_size,
                                                                batch_size=batch_size,
                                                                class_mode='categorical',
                                                                classes= list(map(str,range(0,43))),
                                                                subset='validation') # establece como datos de validación

```

Found 29416 images belonging to 43 classes.
Found 9793 images belonging to 43 classes.

Modelo CNN para clasificación

```
In [5]: #Modelo CNN
def initialize_model():
    model = Sequential()

    ### Primera convolución y MaxPooling
    model.add(Conv2D(32, (3,3), padding = "same", input_shape=image_size+(3,), activation="relu"))
    model.add(Conv2D(32, (3,3), activation="relu"))
    model.add(MaxPool2D(pool_size=(2,2)))
    #model.add(BatchNormalization())
    model.add(Dropout(0.25))

    ### Tercera convolución
    model.add(Conv2D(64, (3,3), padding='same', activation="relu"))
    model.add(Conv2D(64, (3,3), activation="relu"))
    model.add(MaxPool2D(pool_size=(2,2)))
    #model.add(BatchNormalization())
    model.add(Dropout(0.25))

    ### Aplanamiento
    model.add(Flatten())

    ### Una capa "Dense"
    model.add(Dense(392, activation='relu'))
    model.add(Dropout(0.5))

    ### Última capa con 43 salidas
    model.add(Dense(43, activation='softmax')) #Softmax(43)
    return model

def compile_model(model):

    model.compile(loss='categorical_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy', 'Recall', 'Precision'])
    return model
```

```
In [6]: model_keras = initialize_model()
model_keras = compile_model(model_keras)

model_keras.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 32)	896
conv2d_1 (Conv2D)	(None, 26, 26, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
dropout (Dropout)	(None, 13, 13, 32)	0
conv2d_2 (Conv2D)	(None, 13, 13, 64)	18496
conv2d_3 (Conv2D)	(None, 11, 11, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
dropout_1 (Dropout)	(None, 5, 5, 64)	0
flatten (Flatten)	(None, 1600)	0
dense (Dense)	(None, 392)	627592
dropout_2 (Dropout)	(None, 392)	0
dense_1 (Dense)	(None, 43)	16899
=====		
Total params: 710,059		
Trainable params: 710,059		
Non-trainable params: 0		

```
In [7]: # Crea un callback de EarlyStopping que detiene el entrenamiento cuando el accuracy de validación deja de mejorar
early_stop = EarlyStopping(monitor = 'val_accuracy',
                           mode = 'max',
                           patience = 7,
                           restore_best_weights = True,
                           verbose=1)

# Crea un callback de ModelCheckpoint que guarda el modelo con la mejor accuracy de validación
model_checkpoint = ModelCheckpoint('best_model.h5',
                                   monitor='val_accuracy',
                                   mode='max',
```

```
verbose=1,  
save_best_only=True)
```

Entrenamiento de modelo CNN

```
In [8]: %%time  
  
history = model_keras.fit(train_generator,  
                           epochs=101,  
                           validation_data=validation_generator,  
                           callbacks=[early_stop, model_checkpoint])
```

Epoch 1/101
229/230 [=====>.] - ETA: 0s - loss: 2.2237 - accuracy: 0.3723 - recall: 0.2011 - precision: 0.7761
Epoch 1: val_accuracy improved from -inf to 0.65006, saving model to best_model.h5
230/230 [=====] - 17s 60ms/step - loss: 2.2190 - accuracy: 0.3734 - recall: 0.2024 - precision: 0.7761 - val_loss: 1.1673 - val_accuracy: 0.6501 - val_recall: 0.4419 - val_precision: 0.8339
Epoch 2/101
230/230 [=====] - ETA: 0s - loss: 0.7952 - accuracy: 0.7392 - recall: 0.6337 - precision: 0.8521
Epoch 2: val_accuracy improved from 0.65006 to 0.78679, saving model to best_model.h5
230/230 [=====] - 13s 57ms/step - loss: 0.7952 - accuracy: 0.7392 - recall: 0.6337 - precision: 0.8521 - val_loss: 0.6740 - val_accuracy: 0.7868 - val_recall: 0.7178 - val_precision: 0.8693
Epoch 3/101
229/230 [=====>.] - ETA: 0s - loss: 0.4520 - accuracy: 0.8499 - recall: 0.8032 - precision: 0.8981
Epoch 3: val_accuracy improved from 0.78679 to 0.84571, saving model to best_model.h5
230/230 [=====] - 13s 56ms/step - loss: 0.4525 - accuracy: 0.8496 - recall: 0.8030 - precision: 0.8979 - val_loss: 0.5391 - val_accuracy: 0.8457 - val_recall: 0.8028 - val_precision: 0.8942
Epoch 4/101
229/230 [=====>.] - ETA: 0s - loss: 0.3291 - accuracy: 0.8910 - recall: 0.8612 - precision: 0.9219
Epoch 4: val_accuracy improved from 0.84571 to 0.85980, saving model to best_model.h5
230/230 [=====] - 13s 55ms/step - loss: 0.3290 - accuracy: 0.8910 - recall: 0.8612 - precision: 0.9220 - val_loss: 0.5034 - val_accuracy: 0.8598 - val_recall: 0.8250 - val_precision: 0.8999
Epoch 5/101
230/230 [=====] - ETA: 0s - loss: 0.2578 - accuracy: 0.9136 - recall: 0.8942 - precision: 0.9338
Epoch 5: val_accuracy improved from 0.85980 to 0.86878, saving model to best_model.h5
230/230 [=====] - 12s 53ms/step - loss: 0.2578 - accuracy: 0.9136 - recall: 0.8942 - precision: 0.9338 - val_loss: 0.4518 - val_accuracy: 0.8688 - val_recall: 0.8416 - val_precision: 0.9062
Epoch 6/101
230/230 [=====] - ETA: 0s - loss: 0.2201 - accuracy: 0.9263 - recall: 0.9117 - precision: 0.9424
Epoch 6: val_accuracy improved from 0.86878 to 0.88145, saving model to best_model.h5
230/230 [=====] - 12s 51ms/step - loss: 0.2201 - accuracy: 0.9263 - recall: 0.9117 - precision: 0.9424 - val_loss: 0.4409 - val_accuracy: 0.8814 - val_recall: 0.8616 - val_precision: 0.9089
Epoch 7/101
230/230 [=====] - ETA: 0s - loss: 0.1803 - accuracy: 0.9398 - recall: 0.9290 - precision: 0.9520
Epoch 7: val_accuracy improved from 0.88145 to 0.88400, saving model to best_model.h5
230/230 [=====] - 12s 52ms/step - loss: 0.1803 - accuracy: 0.9398 - recall: 0.9290 - precision: 0.9520 - val_loss: 0.4177 - val_accuracy: 0.8840 - val_recall: 0.8669 - val_precision: 0.9110
Epoch 8/101
230/230 [=====] - ETA: 0s - loss: 0.1634 - accuracy: 0.9460 - recall: 0.9364 - precision: 0.9560
Epoch 8: val_accuracy improved from 0.88400 to 0.89053, saving model to best_model.h5
230/230 [=====] - 12s 52ms/step - loss: 0.1634 - accuracy: 0.9460 - recall: 0.9364 - precision: 0.9560 - val_loss: 0.4098 - val_accuracy: 0.8905 - val_recall: 0.8756 - val_precision: 0.9154
Epoch 9/101
229/230 [=====>.] - ETA: 0s - loss: 0.1469 - accuracy: 0.9516 - recall: 0.9434 - precision: 0.9601
Epoch 9: val_accuracy did not improve from 0.89053
230/230 [=====] - 12s 51ms/step - loss: 0.1466 - accuracy: 0.9517 - recall: 0.9435 - precision: 0.9601 - val_loss: 0.4746 - val_accuracy: 0.8867 - val_recall: 0.8739 - val_precision: 0.9055
Epoch 10/101
229/230 [=====>.] - ETA: 0s - loss: 0.1327 - accuracy: 0.9563 - recall: 0.9488 - precision: 0.9640
Epoch 10: val_accuracy improved from 0.89053 to 0.90473, saving model to best_model.h5

230/230 [=====] - 12s 54ms/step - loss: 0.1327 - accuracy: 0.9564 - recall: 0.9489 - precision: 0.9640 - val_loss: 0.3591 - val_accuracy: 0.9047 - val_recall: 0.8915 - val_precision: 0.9267
Epoch 11/101
230/230 [=====] - ETA: 0s - loss: 0.1196 - accuracy: 0.9590 - recall: 0.9536 - precision: 0.9660
Epoch 11: val_accuracy did not improve from 0.90473
230/230 [=====] - 13s 54ms/step - loss: 0.1196 - accuracy: 0.9590 - recall: 0.9536 - precision: 0.9660 - val_loss: 0.3925 - val_accuracy: 0.8998 - val_recall: 0.8896 - val_precision: 0.9182
Epoch 12/101
230/230 [=====] - ETA: 0s - loss: 0.1093 - accuracy: 0.9633 - recall: 0.9584 - precision: 0.9692
Epoch 12: val_accuracy improved from 0.90473 to 0.90851, saving model to best_model.h5
230/230 [=====] - 12s 54ms/step - loss: 0.1093 - accuracy: 0.9633 - recall: 0.9584 - precision: 0.9692 - val_loss: 0.3908 - val_accuracy: 0.9085 - val_recall: 0.9011 - val_precision: 0.9237
Epoch 13/101
229/230 [=====>.] - ETA: 0s - loss: 0.1051 - accuracy: 0.9651 - recall: 0.9611 - precision: 0.9713
Epoch 13: val_accuracy did not improve from 0.90851
230/230 [=====] - 13s 55ms/step - loss: 0.1051 - accuracy: 0.9651 - recall: 0.9611 - precision: 0.9713 - val_loss: 0.4064 - val_accuracy: 0.9015 - val_recall: 0.8907 - val_precision: 0.9212
Epoch 14/101
229/230 [=====>.] - ETA: 0s - loss: 0.0958 - accuracy: 0.9693 - recall: 0.9652 - precision: 0.9737
Epoch 14: val_accuracy improved from 0.90851 to 0.90922, saving model to best_model.h5
230/230 [=====] - 13s 55ms/step - loss: 0.0957 - accuracy: 0.9693 - recall: 0.9653 - precision: 0.9737 - val_loss: 0.3793 - val_accuracy: 0.9092 - val_recall: 0.8986 - val_precision: 0.9252
Epoch 15/101
229/230 [=====>.] - ETA: 0s - loss: 0.0943 - accuracy: 0.9692 - recall: 0.9650 - precision: 0.9734
Epoch 15: val_accuracy improved from 0.90922 to 0.91525, saving model to best_model.h5
230/230 [=====] - 13s 55ms/step - loss: 0.0943 - accuracy: 0.9691 - recall: 0.9649 - precision: 0.9733 - val_loss: 0.3837 - val_accuracy: 0.9152 - val_recall: 0.9093 - val_precision: 0.9272
Epoch 16/101
230/230 [=====] - ETA: 0s - loss: 0.0863 - accuracy: 0.9713 - recall: 0.9677 - precision: 0.9750
Epoch 16: val_accuracy improved from 0.91525 to 0.92137, saving model to best_model.h5
230/230 [=====] - 13s 55ms/step - loss: 0.0863 - accuracy: 0.9713 - recall: 0.9677 - precision: 0.9750 - val_loss: 0.3321 - val_accuracy: 0.9214 - val_recall: 0.9147 - val_precision: 0.9355
Epoch 17/101
229/230 [=====>.] - ETA: 0s - loss: 0.0806 - accuracy: 0.9727 - recall: 0.9691 - precision: 0.9760
Epoch 17: val_accuracy did not improve from 0.92137
230/230 [=====] - 13s 56ms/step - loss: 0.0807 - accuracy: 0.9726 - recall: 0.9691 - precision: 0.9760 - val_loss: 0.3914 - val_accuracy: 0.9102 - val_recall: 0.9028 - val_precision: 0.9239
Epoch 18/101
230/230 [=====] - ETA: 0s - loss: 0.0752 - accuracy: 0.9751 - recall: 0.9725 - precision: 0.9780
Epoch 18: val_accuracy did not improve from 0.92137
230/230 [=====] - 13s 57ms/step - loss: 0.0752 - accuracy: 0.9751 - recall: 0.9725 - precision: 0.9780 - val_loss: 0.3759 - val_accuracy: 0.9120 - val_recall: 0.9046 - val_precision: 0.9274
Epoch 19/101
229/230 [=====>.] - ETA: 0s - loss: 0.0730 - accuracy: 0.9751 - recall: 0.9731 - precision: 0.9783
Epoch 19: val_accuracy did not improve from 0.92137
230/230 [=====] - 13s 57ms/step - loss: 0.0734 - accuracy: 0.9750 - recall: 0.9729 - precision: 0.9782 - val_loss: 0.3318 - val_accuracy: 0.9181 - val_recall: 0.9095 - val_precision: 0.9290
Epoch 20/101

230/230 [=====] - ETA: 0s - loss: 0.0715 - accuracy: 0.9761 - recall: 0.9732 - precision: 0.9795
Epoch 20: val_accuracy did not improve from 0.92137
230/230 [=====] - 13s 57ms/step - loss: 0.0715 - accuracy: 0.9761 - recall: 0.9732 - precision: 0.9795 - val_loss: 0.4076 - val_accuracy: 0.9086 - val_recall: 0.9022 - val_precision: 0.9249
Epoch 21/101
229/230 [=====>.] - ETA: 0s - loss: 0.0683 - accuracy: 0.9770 - recall: 0.9744 - precision: 0.9797
Epoch 21: val_accuracy improved from 0.92137 to 0.92362, saving model to best_model.h5
230/230 [=====] - 13s 57ms/step - loss: 0.0685 - accuracy: 0.9768 - recall: 0.9743 - precision: 0.9796 - val_loss: 0.3607 - val_accuracy: 0.9236 - val_recall: 0.9199 - val_precision: 0.9326
Epoch 22/101
229/230 [=====>.] - ETA: 0s - loss: 0.0647 - accuracy: 0.9779 - recall: 0.9758 - precision: 0.9809
Epoch 22: val_accuracy did not improve from 0.92362
230/230 [=====] - 13s 57ms/step - loss: 0.0648 - accuracy: 0.9778 - recall: 0.9758 - precision: 0.9808 - val_loss: 0.3718 - val_accuracy: 0.9193 - val_recall: 0.9136 - val_precision: 0.9293
Epoch 23/101
230/230 [=====] - ETA: 0s - loss: 0.0618 - accuracy: 0.9790 - recall: 0.9771 - precision: 0.9812
Epoch 23: val_accuracy did not improve from 0.92362
230/230 [=====] - 14s 59ms/step - loss: 0.0618 - accuracy: 0.9790 - recall: 0.9771 - precision: 0.9812 - val_loss: 0.4598 - val_accuracy: 0.9147 - val_recall: 0.9079 - val_precision: 0.9243
Epoch 24/101
230/230 [=====] - ETA: 0s - loss: 0.0578 - accuracy: 0.9805 - recall: 0.9788 - precision: 0.9830
Epoch 24: val_accuracy did not improve from 0.92362
230/230 [=====] - 14s 59ms/step - loss: 0.0578 - accuracy: 0.9805 - recall: 0.9788 - precision: 0.9830 - val_loss: 0.3831 - val_accuracy: 0.9174 - val_recall: 0.9123 - val_precision: 0.9288
Epoch 25/101
229/230 [=====>.] - ETA: 0s - loss: 0.0608 - accuracy: 0.9797 - recall: 0.9776 - precision: 0.9819
Epoch 25: val_accuracy improved from 0.92362 to 0.92495, saving model to best_model.h5
230/230 [=====] - 14s 59ms/step - loss: 0.0609 - accuracy: 0.9795 - recall: 0.9775 - precision: 0.9818 - val_loss: 0.3512 - val_accuracy: 0.9249 - val_recall: 0.9209 - val_precision: 0.9351
Epoch 26/101
229/230 [=====>.] - ETA: 0s - loss: 0.0551 - accuracy: 0.9819 - recall: 0.9803 - precision: 0.9840
Epoch 26: val_accuracy did not improve from 0.92495
230/230 [=====] - 14s 60ms/step - loss: 0.0553 - accuracy: 0.9819 - recall: 0.9803 - precision: 0.9839 - val_loss: 0.4730 - val_accuracy: 0.9122 - val_recall: 0.9080 - val_precision: 0.9221
Epoch 27/101
230/230 [=====] - ETA: 0s - loss: 0.0538 - accuracy: 0.9820 - recall: 0.9802 - precision: 0.9839
Epoch 27: val_accuracy improved from 0.92495 to 0.92944, saving model to best_model.h5
230/230 [=====] - 14s 60ms/step - loss: 0.0538 - accuracy: 0.9820 - recall: 0.9802 - precision: 0.9839 - val_loss: 0.3188 - val_accuracy: 0.9294 - val_recall: 0.9231 - val_precision: 0.9392
Epoch 28/101
230/230 [=====] - ETA: 0s - loss: 0.0565 - accuracy: 0.9802 - recall: 0.9789 - precision: 0.9826
Epoch 28: val_accuracy did not improve from 0.92944
230/230 [=====] - 14s 59ms/step - loss: 0.0565 - accuracy: 0.9802 - recall: 0.9789 - precision: 0.9826 - val_loss: 0.4257 - val_accuracy: 0.9180 - val_recall: 0.9137 - val_precision: 0.9252
Epoch 29/101
230/230 [=====] - ETA: 0s - loss: 0.0534 - accuracy: 0.9825 - recall: 0.9811 - precision: 0.9841
Epoch 29: val_accuracy did not improve from 0.92944
230/230 [=====] - 14s 59ms/step - loss: 0.0534 - accuracy: 0.9825 - recall: 0.9811 - precision: 0.9841 - val_loss:

0.4267 - val_accuracy: 0.9125 - val_recall: 0.9089 - val_precision: 0.9215
Epoch 30/101
229/230 [=====>.] - ETA: 0s - loss: 0.0521 - accuracy: 0.9830 - recall: 0.9816 - precision: 0.9848
Epoch 30: val_accuracy did not improve from 0.92944
230/230 [=====] - 14s 60ms/step - loss: 0.0523 - accuracy: 0.9830 - recall: 0.9816 - precision: 0.9848 - val_loss: 0.3900 - val_accuracy: 0.9261 - val_recall: 0.9236 - val_precision: 0.9323
Epoch 31/101
230/230 [=====] - ETA: 0s - loss: 0.0461 - accuracy: 0.9842 - recall: 0.9831 - precision: 0.9860
Epoch 31: val_accuracy did not improve from 0.92944
230/230 [=====] - 14s 60ms/step - loss: 0.0461 - accuracy: 0.9842 - recall: 0.9831 - precision: 0.9860 - val_loss: 0.3319 - val_accuracy: 0.9284 - val_recall: 0.9233 - val_precision: 0.9362
Epoch 32/101
230/230 [=====] - ETA: 0s - loss: 0.0524 - accuracy: 0.9833 - recall: 0.9821 - precision: 0.9852
Epoch 32: val_accuracy improved from 0.92944 to 0.93414, saving model to best_model.h5
230/230 [=====] - 14s 60ms/step - loss: 0.0524 - accuracy: 0.9833 - recall: 0.9821 - precision: 0.9852 - val_loss: 0.2969 - val_accuracy: 0.9341 - val_recall: 0.9291 - val_precision: 0.9407
Epoch 33/101
229/230 [=====>.] - ETA: 0s - loss: 0.0510 - accuracy: 0.9836 - recall: 0.9823 - precision: 0.9849
Epoch 33: val_accuracy did not improve from 0.93414
230/230 [=====] - 14s 60ms/step - loss: 0.0509 - accuracy: 0.9837 - recall: 0.9824 - precision: 0.9850 - val_loss: 0.3916 - val_accuracy: 0.9227 - val_recall: 0.9181 - val_precision: 0.9321
Epoch 34/101
230/230 [=====] - ETA: 0s - loss: 0.0506 - accuracy: 0.9835 - recall: 0.9822 - precision: 0.9854
Epoch 34: val_accuracy did not improve from 0.93414
230/230 [=====] - 15s 64ms/step - loss: 0.0506 - accuracy: 0.9835 - recall: 0.9822 - precision: 0.9854 - val_loss: 0.3737 - val_accuracy: 0.9253 - val_recall: 0.9213 - val_precision: 0.9318
Epoch 35/101
230/230 [=====] - ETA: 0s - loss: 0.0421 - accuracy: 0.9856 - recall: 0.9845 - precision: 0.9870
Epoch 35: val_accuracy did not improve from 0.93414
230/230 [=====] - 14s 63ms/step - loss: 0.0421 - accuracy: 0.9856 - recall: 0.9845 - precision: 0.9870 - val_loss: 0.4074 - val_accuracy: 0.9254 - val_recall: 0.9217 - val_precision: 0.9311
Epoch 36/101
230/230 [=====] - ETA: 0s - loss: 0.0459 - accuracy: 0.9848 - recall: 0.9835 - precision: 0.9858
Epoch 36: val_accuracy did not improve from 0.93414
230/230 [=====] - 14s 62ms/step - loss: 0.0459 - accuracy: 0.9848 - recall: 0.9835 - precision: 0.9858 - val_loss: 0.4383 - val_accuracy: 0.9122 - val_recall: 0.9069 - val_precision: 0.9239
Epoch 37/101
230/230 [=====] - ETA: 0s - loss: 0.0441 - accuracy: 0.9849 - recall: 0.9840 - precision: 0.9862
Epoch 37: val_accuracy did not improve from 0.93414
230/230 [=====] - 14s 62ms/step - loss: 0.0441 - accuracy: 0.9849 - recall: 0.9840 - precision: 0.9862 - val_loss: 0.3715 - val_accuracy: 0.9262 - val_recall: 0.9227 - val_precision: 0.9331
Epoch 38/101
230/230 [=====] - ETA: 0s - loss: 0.0435 - accuracy: 0.9851 - recall: 0.9841 - precision: 0.9869
Epoch 38: val_accuracy did not improve from 0.93414
230/230 [=====] - 14s 61ms/step - loss: 0.0435 - accuracy: 0.9851 - recall: 0.9841 - precision: 0.9869 - val_loss: 0.3656 - val_accuracy: 0.9305 - val_recall: 0.9277 - val_precision: 0.9372
Epoch 39/101
229/230 [=====>.] - ETA: 0s - loss: 0.0434 - accuracy: 0.9857 - recall: 0.9848 - precision: 0.9869Restoring model we

ights from the end of the best epoch: 32.

Epoch 39: val_accuracy did not improve from 0.93414

230/230 [=====] - 13s 56ms/step - loss: 0.0435 - accuracy: 0.9856 - recall: 0.9848 - precision: 0.9869 - val_loss: 0.4024 - val_accuracy: 0.9252 - val_recall: 0.9223 - val_precision: 0.9323

Epoch 39: early stopping

CPU times: total: 12min 40s

Wall time: 8min 38s



Resultados de entrenamiento

```
In [9]: # Plot the loss and accuracy curves
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend(['Train', 'Validation'])

plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(['Train', 'Validation'])

plt.tight_layout()
plt.show()
```



Evaluación de modelo de clasificación

```
In [6]: # Cargar la tabla pandas con las rutas de las imágenes y las clases
df = pd.read_csv("data/Test.csv")

# Definir una función que carga una imagen desde una ruta y la preprocesa
def load_image(path):
    # Leer el archivo de imagen
    image = tf.io.read_file(path)
    # Decodificar la imagen como JPEG
    image = tf.image.decode_png(image, channels=3)
    # Cambiar el tamaño de la imagen a un tamaño fijo
    image = tf.image.resize(image, image_size)
    # Normalizar los valores de píxeles al rango [0, 1]
    image = image / 255.0
    return image

# Definir una función que carga un lote de imágenes y etiquetas desde la tabla pandas
def load_batch(df, batch_size):
    # Barajar la tabla
    df = df.sample(frac=1).reset_index(drop=True)
    # Recorrer la tabla en lotes
    for i in range(0, len(df), batch_size):
        # Obtener el lote de rutas y clases
        batch_paths = df["Path"][i:i+batch_size]
        batch_classes = df["ClassId"][i:i+batch_size]
```

```

# Cargar el lote de imágenes
batch_images = [load_image('data/' + path) for path in batch_paths]
# Convertir el lote de imágenes y etiquetas en tensores
batch_images = tf.stack(batch_images)
# Convertir el lote de etiquetas en tensores categóricos
batch_labels = tf.keras.utils.to_categorical(batch_classes, num_classes=43)
yield batch_images, batch_labels

# Definir una función de cargador de datos que usa la función load_batch
def data_loader(df, batch_size):
    # Crear un conjunto de datos a partir de la función load_batch
    dataset = tf.data.Dataset.from_generator(
        lambda: load_batch(df, batch_size),
        output_types=(tf.float32, tf.float32),
        output_shapes=(None, image_size[0], image_size[1], 3), [None, None])
    )
    return dataset

# Crear un cargador de datos con un tamaño de lote de 32
data_loader = data_loader(df, 32)

```

```

In [7]: # Cargar el mejor modelo
model = tf.keras.models.load_model("best_model.h5")

# Evaluar el modelo con los datos de Test
model.evaluate(data_loader)

```

395/395 [=====] - 113s 277ms/step - loss: 0.2606 - accuracy: 0.9435 - recall: 0.9408 - precision: 0.9494

```
Out[7]: [0.2606205344200134, 0.943467915058136, 0.9407759308815002, 0.949420690536499]
```

Modelo para detección de señales en el video

Para complementar la clasificación de las señales de tránsito, se utilizó el set de datos de GTSDb (German Traffic Sign Detection Benchmark) el cual es una evaluación de detección de imágenes individuales para investigadores con interés en el campo de la visión por computadora, el reconocimiento de patrones y la asistencia al conductor basada en imágenes. Se supone que se presentará en la IEEE International Joint Conference on Neural Networks 2013:

https://benchmark.ini.rub.de/gtsdb_news.html

Datos de entrenamiento

```
In [5]: path_dir="data/YOLO_data/FullIJCNN2013/"

data = pd.read_csv(path_dir+'gt.txt',sep=';',names=['path','left','top','right','bottom','id'])
data.head()
```

```
Out[5]:
```

	path	left	top	right	bottom	id
0	00000.ppm	774	411	815	446	11
1	00001.ppm	983	388	1024	432	40
2	00001.ppm	386	494	442	552	38
3	00001.ppm	973	335	1031	390	13
4	00002.ppm	892	476	1006	592	39

```
In [7]: from PIL import Image
#opening image and converting it into numpy array and checking the size of array.
img = Image.open(path_dir+data['path'][0])
img=np.array(img)
print(img.shape)

img = Image.open(path_dir+data['path'][1])
img=np.array(img)
plt.axis('off')
plt.imshow(img[:,:,:])
plt.show()
```

(800, 1360, 3)



```
In [8]: #copiando la columna id de los datos en df
df = data['id'].copy()
```

```
data['Object Name'] = data['id']
#asignando una sola etiqueta para señal de trafico
for i in range(len(df)):
    if(df[i] in range(0,43)):
        df.loc[i]=0
        data['Object Name'].loc[i]='trafficsign'
    else:
        df.loc[i]=-1
data.head()
```

C:\Users\fgarcia24\AppData\Local\Temp\ipykernel_44008\1312309559.py:8: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
data['Object Name'].loc[i]='trafficsign'

Out[8]:

	path	left	top	right	bottom	id	Object Name
0	00000.ppm	774	411	815	446	11	trafficsign
1	00001.ppm	983	388	1024	432	40	trafficsign
2	00001.ppm	386	494	442	552	38	trafficsign
3	00001.ppm	973	335	1031	390	13	trafficsign
4	00002.ppm	892	476	1006	592	39	trafficsign

Ajustes de dato para el entrenamiento

In [9]:

```
import cv2

#Remover la extension en la columna de path
data['path'] = data['path'].str.slice(0, -4)

# Usando os.walk para recorrer todos los directorios
# y archivos en ellos desde el directorio actual
path_dir="data/YOLO_data/FullIJCNN2013/"
final_df=pd.DataFrame(columns=data.columns)
df_list = [] # crear una lista vacía de DataFrames
for current_dir, dirs, files in os.walk(path_dir):
    for f in files:
        #print(f)
        #Comprobar si un archivo termina con formato .ppm o no
        if f.endswith('.ppm'):
            image_name=f[:-4]
            img=cv2.imread(path_dir + f)
```

```

single_yolo_dat = data.loc[data['path'] == image_name].copy()
#print(single_yolo_dat)
# y, de esta manera, el dataframe inicial no se cambiará
# Comprobando si no hay ninguna anotación para la imagen actual
if single_yolo_dat.isnull().values.all():
    # Eliminando esta imagen de los datos de entrenamiento
    # print(f)
    os.remove(path_dir + f)

#Ahora guardar el resulted_frame en una carpeta dentro de path_dir
else:
    df_list.append(single_yolo_dat) # añadir a la lista de DataFrames
    #Ahora escribiendo y guardando la imagen de formato ppm a formato jpg usando OpenCV
    save_path = 'data/YOLO_data/jpeg_files/' + image_name + '.jpg'
    cv2.imwrite(save_path, img)

final_df = pd.concat(df_list) # concatenar la lista de DataFrames
final_df = final_df[~final_df.index.duplicated(keep='first')]
final_df.sort_index(inplace=True)

```

División del dato para entrenamiento y validación

```

In [11]: from sklearn.model_selection import train_test_split

train, test = train_test_split(final_df['path'].unique(), test_size=0.2, random_state=37)
train.shape, test.shape

```

```

Out[11]: ((592,), (149,))

```

Creación de etiquetas para entrenamiento y validación

```

In [ ]: labels_path = 'data/YOLO_data/labels/'

# Crear una carpeta para almacenar los archivos de etiquetas
os.makedirs(labels_path, exist_ok=True)
os.makedirs(labels_path+'train/', exist_ok=True)
os.makedirs(labels_path+'test/', exist_ok=True)

# Recorrer cada imagen en la tabla
for path, group in final_df.groupby('path'):
    # Obtener el nombre de la imagen sin extensión
    if path in train:
        name = labels_path+'train/'+path+'.txt'
    elif path in test:

```

```

name = labels_path+'test/'+path+'.txt'
# Abrir un archivo de texto con el mismo nombre que la imagen
with open(name, 'w') as f:
    # Recorrer cada objeto en la imagen
    for row in group.itertuples():
        # Obtener el id de clase del objeto, las coordenadas izquierda, superior, derecha e inferior
        class_id = 0 #row.id # modificado para detectar cualquier señal de tráfico
        left = row.left
        top = row.top
        right = row.right
        bottom = row.bottom
        # Calcular el centro normalizado x, y, ancho y alto del cuadro delimitador
        x = round((left + right) / 2 / 1360,2) # ancho de la imagen
        y = round((top + bottom) / 2 / 800,2) # altura de la imagen
        w = round((right - left) / 1360,2)
        h = round((bottom - top) / 800,2)

        # Escribir los datos de la etiqueta en el archivo de texto en formato YOLOv8
        f.write(f'{class_id} {x} {y} {w} {h}\n')

```

División de las imágenes de acuerdo a las listas de entrenamiento y validación

In [13]:

```

import shutil

jpeg_files = 'data/YOLO_data/jpeg_files/'
images_path = 'data/YOLO_data/images/'

# Crear las carpetas de entrenamiento y prueba si no existen
os.makedirs(images_path, exist_ok=True)
os.makedirs(images_path+'train/', exist_ok=True)
os.makedirs(images_path+'test/', exist_ok=True)

# Recorrer el array de entrenamiento y copiar las imágenes jpeg a la carpeta de entrenamiento
for image in train:
    # Obtener la ruta completa del archivo de imagen
    src = os.path.join(jpeg_files, image + '.jpg')
    # Copiar el archivo de imagen a la carpeta de entrenamiento
    shutil.copy(src, images_path+'train')

# Recorrer el array de prueba y copiar las imágenes jpeg a la carpeta de prueba
for image in test:
    # Obtener la ruta completa del archivo de imagen
    src = os.path.join(jpeg_files, image + '.jpg')
    # Copiar el archivo de imagen a la carpeta de prueba
    shutil.copy(src, images_path+'test')

```


Entrenamiento de modelo YOLO

```
In [14]: import torch
from ultralytics import YOLO

os.environ["KMP_DUPLICATE_LIB_OK"]="TRUE"
```

```
In [15]: model = YOLO('yolov8s.yaml').load('yolov8s.pt')
```

	from	n	params	module	arguments
0	-1	1	928	ultralytics.nn.modules.conv.Conv	[3, 32, 3, 2]
1	-1	1	18560	ultralytics.nn.modules.conv.Conv	[32, 64, 3, 2]
2	-1	1	29056	ultralytics.nn.modules.block.C2f	[64, 64, 1, True]
3	-1	1	73984	ultralytics.nn.modules.conv.Conv	[64, 128, 3, 2]
4	-1	2	197632	ultralytics.nn.modules.block.C2f	[128, 128, 2, True]
5	-1	1	295424	ultralytics.nn.modules.conv.Conv	[128, 256, 3, 2]
6	-1	2	788480	ultralytics.nn.modules.block.C2f	[256, 256, 2, True]
7	-1	1	1180672	ultralytics.nn.modules.conv.Conv	[256, 512, 3, 2]
8	-1	1	1838080	ultralytics.nn.modules.block.C2f	[512, 512, 1, True]
9	-1	1	656896	ultralytics.nn.modules.block.SPPF	[512, 512, 5]
10	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
11	[-1, 6]	1	0	ultralytics.nn.modules.conv.Concat	[1]
12	-1	1	591360	ultralytics.nn.modules.block.C2f	[768, 256, 1]
13	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
14	[-1, 4]	1	0	ultralytics.nn.modules.conv.Concat	[1]
15	-1	1	148224	ultralytics.nn.modules.block.C2f	[384, 128, 1]
16	-1	1	147712	ultralytics.nn.modules.conv.Conv	[128, 128, 3, 2]
17	[-1, 12]	1	0	ultralytics.nn.modules.conv.Concat	[1]
18	-1	1	493056	ultralytics.nn.modules.block.C2f	[384, 256, 1]
19	-1	1	590336	ultralytics.nn.modules.conv.Conv	[256, 256, 3, 2]
20	[-1, 9]	1	0	ultralytics.nn.modules.conv.Concat	[1]
21	-1	1	1969152	ultralytics.nn.modules.block.C2f	[768, 512, 1]
22	[15, 18, 21]	1	2147008	ultralytics.nn.modules.head.Detect	[80, [128, 256, 512]]

YOLOv8s summary: 225 layers, 11166560 parameters, 11166544 gradients

Transferred 355/355 items from pretrained weights

```
In [16]: %%time
results = model.train(data='yolo-traffic signs.yaml', epochs=83, imgsz=640, verbose=True, patience=13, single_cls=True,
                      conf=0.5, iou=0.7,
                      project='traffic signs_yolo8s_640', name='train', pretrained=True, optimizer='SGD',
                      dropout=0.1, cls=0.5,
                      device=0, workers=16, batch=32, save_conf=True,
                      augment=True, mosaic=False, mixup=False,
                      degrees=7, flipud=0, scale=0.1,
```

```
hsv_v=0.75, hsv_h=0.085, hsv_s=0.75  
)
```

Ultralytics YOLOv8.0.106 Python-3.9.16 torch-2.0.1+cu117 CUDA:0 (NVIDIA RTX A4500 Laptop GPU, 16384MiB)

yolo\engine\trainer: task=detect, mode=train, model=yolov8s.yaml, data=yolo-traffic signs.yaml, epochs=83, patience=13, batch=32, imgsz=640, save=True, save_period=-1, cache=False, device=0, workers=16, project=traffic signs_yolo8s_640, name=train, exist_ok=False, pretrained=True, optimizer=SGD, verbose=True, seed=0, deterministic=True, single_cls=True, rect=False, cos_lr=False, close_mosaic=0, resume=False, amp=True, overlap_mask=True, mask_ratio=4, dropout=0.1, val=True, split=val, save_json=False, save_hybrid=False, conf=0.5, iou=0.7, max_det=300, half=False, dnn=False, plots=True, source=None, show=False, save_txt=False, save_conf=True, save_crop=False, show_labels=True, show_conf=True, vid_stride=1, line_width=None, visualize=False, augment=True, agnostic_nms=False, classes=None, retina_masks=False, boxes=True, format=torchscript, keras=False, optimize=False, int8=False, dynamic=False, simplify=False, opset=None, workspace=4, nms=False, lr0=0.01, lrf=0.01, momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=7.5, cls=0.5, dfl=1.5, pose=12.0, kobj=1.0, label_smoothing=0.0, nbs=64, hsv_h=0.085, hsv_s=0.75, hsv_v=0.75, degrees=7, translate=0.1, scale=0.1, shear=0.0, perspective=0.0, flipud=0, fliplr=0.5, mosaic=False, mixup=False, copy_paste=0.0, cfg=None, v5loader=False, tracker=botsort.yaml, save_dir=traffic signs_yolo8s_640\train2

Overriding model.yaml nc=80 with nc=1

	from	n	params	module	arguments
0	-1	1	928	ultralytics.nn.modules.conv.Conv	[3, 32, 3, 2]
1	-1	1	18560	ultralytics.nn.modules.conv.Conv	[32, 64, 3, 2]
2	-1	1	29056	ultralytics.nn.modules.block.C2f	[64, 64, 1, True]
3	-1	1	73984	ultralytics.nn.modules.conv.Conv	[64, 128, 3, 2]
4	-1	2	197632	ultralytics.nn.modules.block.C2f	[128, 128, 2, True]
5	-1	1	295424	ultralytics.nn.modules.conv.Conv	[128, 256, 3, 2]
6	-1	2	788480	ultralytics.nn.modules.block.C2f	[256, 256, 2, True]
7	-1	1	1180672	ultralytics.nn.modules.conv.Conv	[256, 512, 3, 2]
8	-1	1	1838080	ultralytics.nn.modules.block.C2f	[512, 512, 1, True]
9	-1	1	656896	ultralytics.nn.modules.block.SPPF	[512, 512, 5]
10	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
11	[-1, 6]	1	0	ultralytics.nn.modules.conv.Concat	[1]
12	-1	1	591360	ultralytics.nn.modules.block.C2f	[768, 256, 1]
13	-1	1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
14	[-1, 4]	1	0	ultralytics.nn.modules.conv.Concat	[1]
15	-1	1	148224	ultralytics.nn.modules.block.C2f	[384, 128, 1]
16	-1	1	147712	ultralytics.nn.modules.conv.Conv	[128, 128, 3, 2]
17	[-1, 12]	1	0	ultralytics.nn.modules.conv.Concat	[1]
18	-1	1	493056	ultralytics.nn.modules.block.C2f	[384, 256, 1]
19	-1	1	590336	ultralytics.nn.modules.conv.Conv	[256, 256, 3, 2]
20	[-1, 9]	1	0	ultralytics.nn.modules.conv.Concat	[1]
21	-1	1	1969152	ultralytics.nn.modules.block.C2f	[768, 512, 1]
22	[15, 18, 21]	1	2116435	ultralytics.nn.modules.head.Detect	[1, [128, 256, 512]]

YOLOv8s summary: 225 layers, 11135987 parameters, 11135971 gradients


Transferred 349/355 items from pretrained weights

TensorBoard: Start with 'tensorboard --logdir traffic signs_yolo8s_640\train2', view at <http://localhost:6006/>

AMP: running Automatic Mixed Precision (AMP) checks with YOLOv8n...

AMP: checks passed

optimizer: SGD(lr=0.01) with parameter groups 57 weight(decay=0.0), 64 weight(decay=0.0005), 63 bias

train: Scanning C:\Users\fgarcia24\OneDrive - Schlumberger\Python\fco-parga\mna\mna-navegacion_autonoma\Actividad-4.2\data\YOLO_data\labels\train.cache... 592 images, 0 backgrounds, 0 corrupt: 100% 592/592 [00:00<?, ?it/s]

```

train: WARNING C:\Users\fgarcia24\OneDrive - Schlumberger\Python\fco-parga\mna\mna-navegacion_autonoma\Actividad-4.2\data\YOLO_data\images
\train\00340.jpg: 1 duplicate labels removed
val: Scanning C:\Users\fgarcia24\OneDrive - Schlumberger\Python\fco-parga\mna\mna-navegacion_autonoma\Actividad-4.2\data\YOLO_data\labels\t
est.cache... 149 images, 0 backgrounds, 0 corrupt: 100%[██████████] 149/149 [00:00<?, ?it/s]
Plotting labels to trafficsigns_yolo8s_640\train2\labels.jpg...
Image sizes 640 train, 640 val
Using 16 dataloader workers
Logging results to trafficsigns_yolo8s_640\train2
Starting training for 83 epochs...

```

Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
1/83	7.47G	2.219	14.68	1.271	20	640: 100%[██████████] 19/19 [00:17<00:00, 1.10it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100%[██████████] 3/3 [00:00<00:00, 3.99it/s]
	all	149	243	0.87	0.276	0.576 0.326
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
2/83	7.69G	1.717	1.345	1.071	28	640: 100%[██████████] 19/19 [00:05<00:00, 3.74it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100%[██████████] 3/3 [00:00<00:00, 4.14it/s]
	all	149	243	0.896	0.354	0.632 0.374
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
3/83	7.53G	1.695	1.195	1.028	23	640: 100%[██████████] 19/19 [00:05<00:00, 3.66it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100%[██████████] 3/3 [00:00<00:00, 3.36it/s]
	all	149	243	0.835	0.708	0.803 0.415
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
4/83	7.55G	1.641	1.215	1.037	27	640: 100%[██████████] 19/19 [00:05<00:00, 3.73it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100%[██████████] 3/3 [00:00<00:00, 3.02it/s]
	all	149	243	0.818	0.724	0.817 0.378
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
5/83	7.53G	1.688	1.145	1.032	28	640: 100%[██████████] 19/19 [00:05<00:00, 3.72it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100%[██████████] 3/3 [00:01<00:00, 2.99it/s]
	all	149	243	0.836	0.758	0.838 0.407
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
6/83	7.55G	1.749	1.162	1.075	23	640: 100%[██████████] 19/19 [00:05<00:00, 3.72it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100%[██████████] 3/3 [00:00<00:00, 3.53it/s]
	all	149	243	0.773	0.564	0.682 0.312
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
7/83	7.68G	1.717	1.133	1.041	30	640: 100%[██████████] 19/19 [00:05<00:00, 3.68it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100%[██████████] 3/3 [00:00<00:00, 3.13it/s]
	all	149	243	0.655	0.658	0.688 0.289
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
8/83	7.55G	1.703	1.028	1.042	25	640: 100%[██████████] 19/19 [00:05<00:00, 3.67it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100%[██████████] 3/3 [00:00<00:00, 3.26it/s]

	all	149	243	0.851	0.675	0.788	0.352
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
9/83	7.53G	1.732	0.9913	1.051	29	640: 100%	19/19 [00:05<00:00, 3.66it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
	all	149	243	0.725	0.75	0.761	0.356
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
10/83	7.55G	1.686	1.026	1.028	24	640: 100%	19/19 [00:05<00:00, 3.64it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
	all	149	243	0.737	0.668	0.721	0.361
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
11/83	7.52G	1.715	0.9607	1.03	25	640: 100%	19/19 [00:05<00:00, 3.65it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
	all	149	243	0.934	0.638	0.777	0.396
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
12/83	7.55G	1.699	0.9815	1.038	33	640: 100%	19/19 [00:05<00:00, 3.40it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
	all	149	243	0.823	0.671	0.76	0.359
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
13/83	7.52G	1.679	0.9558	1.027	28	640: 100%	19/19 [00:05<00:00, 3.40it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
	all	149	243	0.792	0.601	0.703	0.348
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
14/83	7.71G	1.69	0.9369	1.041	26	640: 100%	19/19 [00:05<00:00, 3.30it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
	all	149	243	0.897	0.679	0.796	0.391
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
15/83	7.52G	1.659	0.8911	1.037	30	640: 100%	19/19 [00:05<00:00, 3.26it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
	all	149	243	0.883	0.745	0.806	0.385
Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
16/83	7.72G	1.651	0.901	1.03	25	640: 100%	19/19 [00:05<00:00, 3.18it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%
	all	149	243	0.935	0.593	0.763	0.371

Stopping training early as no improvement observed in last 13 epochs. Best results observed at epoch 3, best model saved as best.pt.
To update EarlyStopping(patience=13) pass a new patience value, i.e. `patience=300` or use `patience=0` to disable EarlyStopping.

16 epochs completed in 0.034 hours.

Optimizer stripped from traffic signs_yolo8s_640\train2\weights\last.pt, 22.5MB

Optimizer stripped from traffic signs_yolo8s_640\train2\weights\best.pt, 22.5MB

```
Validating trafficsigns_yolo8s_640\train2\weights\best.pt...
```

```
Ultralytics YOLOv8.0.106 Python-3.9.16 torch-2.0.1+cu117 CUDA:0 (NVIDIA RTX A4500 Laptop GPU, 16384MiB)
```

```
YOLOv8s summary (fused): 168 layers, 11125971 parameters, 0 gradients
```

Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100%	3/3 [00:01<00:00, 2.03it/s]
all	149	243	0.835	0.708	0.803	0.416	

```
Speed: 0.7ms preprocess, 2.7ms inference, 0.0ms loss, 0.8ms postprocess per image
```

```
Results saved to trafficsigns_yolo8s_640\train2
```

```
CPU times: total: 2min 15s
```

```
Wall time: 8min 49s
```

```
In [20]: road_img = cv2.imread('trafficsigns_yolo8s_640/train2/val_batch1_pred.jpg')
```

```
plt.figure(figsize=(20,11))
plt.imshow(road_img)
plt.axis('off')
plt.show()
```




Detector de senales de trafico integrado

```
In [21]: classes = {
    0:'Speed limit (20km/h)',
    1:'Speed limit (30km/h)',
    2:'Speed limit (50km/h)',
    3:'Speed limit (60km/h)',
    4:'Speed limit (70km/h)',
```

```

5:'Speed limit (80km/h)',
6:'End of speed limit (80km/h)',
7:'Speed limit (100km/h)',
8:'Speed limit (120km/h)',
9:'No passing',
10:'No passing veh over 3.5 tons',
11:'Right-of-way at intersection',
12:'Priority road',
13:'Yield',
14:'Stop',
15:'No vehicles',
16:'Veh > 3.5 tons prohibited',
17:'No entry',
18:'General caution',
19:'Dangerous curve left',
20:'Dangerous curve right',
21:'Double curve',
22:'Bumpy road',
23:'Slippery road',
24:'Road narrows on the right',
25:'Road work',
26:'Traffic signals',
27:'Pedestrians',
28:'Children crossing',
29:'Bicycles crossing',
30:'Beware of ice/snow',
31:'Wild animals crossing',
32:'End speed + passing limits',
33:'Turn right ahead',
34:'Turn left ahead',
35:'Ahead only',
36:'Go straight or right',
37:'Go straight or left',
38:'Keep right',
39:'Keep left',
40:'Roundabout mandatory',
41:'End of no passing',
42:'End no passing veh > 3.5 tons'
}

```

```

In [22]: from tqdm.auto import tqdm
from PIL import Image
import tensorflow as tf

class TrafficSignDetector:
    def __init__(self, model_detec, model_classify):
        self.model_detec = model_detec

```



```

self.model_classify = model_classify

def process_frame(self, img):
    # Definir Los parámetros para detección de señales de tráfico
    passed_frame = Image.fromarray(cv2.cvtColor(img, cv2.COLOR_BGR2RGB))
    #Detección de señales de tránsito con YOLO
    detected_trafficsigns = self.model_detec.predict(source=passed_frame, save=False, conf=0.1, verbose=False)
    boxes = detected_trafficsigns[0].boxes.xyxy.cpu().data.numpy().astype(int)

    # Definir una función para redimensionar y normalizar una imagen
    def preprocess_image(image):
        image = tf.image.resize(image, (28, 28))
        image = tf.cast(image, tf.float32)
        image = image/255
        return image

    #Función para crear una caja cuadrada en la señal detectada
    def get_squered_box(box):
        extend = 1
        w = box[2]-box[0]
        h = box[3]-box[1]

        long = (max(w,h)/2)+extend

        center = (np.mean([box[2],box[0]]).astype(int),np.mean([box[3],box[1]]).astype(int))

        square_box = np.array([center[0]-long, center[1]-long,
                               center[0]+long, center[1]+long]).astype(int)
        return square_box, long

    # Clasificación de cada señal detectada dentro el frame
    if boxes.shape[0] > 0:
        for sign in boxes:
            square_box, long = get_squered_box(sign)
            cropped_sign = img[square_box[1]:square_box[3],square_box[0]:square_box[2]]

            # Convertir la región de interés en un tensor de TensorFlow y normalizar
            roi_tensor = preprocess_image(np.array(cropped_sign))
            roi_tensor = tf.expand_dims(roi_tensor, axis=0)

            # Predecir qué tipo de señal de tránsito fue encontrada en la región de interés
            self.prediction = self.model_classify.predict(roi_tensor, verbose=0)
            self.traffic_class = classes.get(np.argmax(self.prediction))

            # Si se detecta una señal de tránsito, dibujar un rectángulo alrededor de la región de interés
            if tf.reduce_any(tf.greater_equal(self.prediction, 0.9)):
                cv2.rectangle(img, (sign[0], sign[1]), (sign[2], sign[3]), (0, 255, 0), 2)
                cv2.putText(img, self.traffic_class, (sign[0], sign[1]-10), cv2.FONT_HERSHEY_SIMPLEX, 0.9, (0,255,0), 2)

```

```

        #print('Traficsignal')

    return img

def process_video(self, input_video_path, output_video_path, start_at=3600, fracction_to_process=0.1):
    frame_set_no=start_at
    cap = cv2.VideoCapture(input_video_path)
    cap.set(cv2.CAP_PROP_POS_FRAMES, frame_set_no)

    width = int(cap.get(cv2.CAP_PROP_FRAME_WIDTH))
    height = int(cap.get(cv2.CAP_PROP_FRAME_HEIGHT))
    fps = int(cap.get(cv2.CAP_PROP_FPS))
    total_frames = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))

    # Calcular el número de fotogramas a procesar (el 10% de los fotogramas totales)
    num_frames_to_process = int(total_frames * fracction_to_process)

    fourcc = cv2.VideoWriter_fourcc(*'mp4v')
    out = cv2.VideoWriter(output_video_path, fourcc, fps, (width, height))

    pbar = tqdm(total=num_frames_to_process, ncols=80, bar_format='{l_bar}{bar}| {n_fmt}/{total_fmt}',
                position=0, leave=True)

    frame_count = 0

    while cap.isOpened():
        ret, frame = cap.read()
        if ret:
            try:
                # Procesar el fotograma con el detector de señales de tráfico
                result_frame = self.process_frame(frame)
                # Escribir el fotograma procesado en el archivo de salida
                out.write(result_frame)
                #cv2.imshow('Processed Frame', result_frame)

                pbar.update(1)
                frame_count += 1

                if frame_count >= num_frames_to_process:
                    break

                #if cv2.waitKey(1) & 0xFF == ord('q'):
                #    break
            except:
                pass
        else:
            break

```

```
pbar.close()

cap.release()
out.release()
#cv2.destroyAllWindows()
```

In [23]: `from ultralytics import YOLO`

```
input_video_path = 'dash_cam.mp4'
output_video_path = 'dash_cams_cnn-Keras.mp4'

model_detec = YOLO('trafficsigns_yolo8s_640/train/weights/best.pt')
model_classify = tf.keras.models.load_model("best_model.h5")
```

In [24]: `%%time`

```
# Crear una instancia de PedestrianDetector con svc_model como el modelo a usar
traffic_sign_detector = TrafficSignDetector(model_detec, model_classify)

# Procesar el archivo de vídeo de entrada y guardar el archivo de vídeo de salida usando el método process_video
traffic_sign_detector.process_video(input_video_path, output_video_path, start_at=3600, fracction_to_process=0.2)
```

0%| | 0/3990

CPU times: total: 7min 29s

Wall time: 6min 24s

Resultado de video procesado

El resultado procesado puede encontrarse en el link de youtube:

<https://youtu.be/ErCkWxos-El>